

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

↗ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
1 cd '/content/drive/MyDrive/Colab Notebooks/'
```

↗ /content/drive/MyDrive/Colab Notebooks

```
1 # Import necessary libraries
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from sklearn.preprocessing import StandardScaler, LabelEncoder
5 from tensorflow.keras.utils import to_categorical
6 import numpy as np
7 !pip install missingno
8 !pip install plotly.express
9 import missingno as msno
10 import matplotlib.pyplot as plt
11 import plotly.express as px
12 from sklearn.compose import ColumnTransformer
13 from sklearn.pipeline import Pipeline
14 from sklearn.impute import SimpleImputer
15 from keras.utils import plot_model
16 from IPython.display import Image
17 from sklearn.metrics import classification_report, accuracy_score
18
19 from keras_tuner.engine.hyperparameters import HyperParameters
20 from keras_tuner import RandomSearch, Hyperband, BayesianOptimization
21 from tensorflow.keras.utils import plot_model
22 from tensorflow.keras.callbacks import LearningRateScheduler, TensorBoard
23 import keras_tuner as kt
24 from sklearn.metrics import classification_report, confusion_matrix
25
```

↗ Requirement already satisfied: missingno in /usr/local/lib/python3.10/dist-packages (0.5.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from missingno) (1.26.4)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from missingno) (3.7.1)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from missingno) (1.13.1)
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (from missingno) (0.13.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (1.3.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (4.53.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (1.4.7)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (24.1)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (3.1.4)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (2.8.2)
Requirement already satisfied: pandas>=1.2 in /usr/local/lib/python3.10/dist-packages (from seaborn->missingno) (2.1.4)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.2->seaborn->missingno) (2024.2)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.2->seaborn->missingno) (2024.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib->missingno) (1.16.0)
Requirement already satisfied: plotly.express in /usr/local/lib/python3.10/dist-packages (0.4.1)
Requirement already satisfied: pandas>=0.20.0 in /usr/local/lib/python3.10/dist-packages (from plotly.express) (2.1.4)
Requirement already satisfied: plotly>=4.1.0 in /usr/local/lib/python3.10/dist-packages (from plotly.express) (5.24.1)
Requirement already satisfied: statsmodels>=0.9.0 in /usr/local/lib/python3.10/dist-packages (from plotly.express) (0.14.3)
Requirement already satisfied: scipy>=0.18 in /usr/local/lib/python3.10/dist-packages (from plotly.express) (1.13.1)
Requirement already satisfied: patsy>=0.5 in /usr/local/lib/python3.10/dist-packages (from plotly.express) (0.5.6)
Requirement already satisfied: numpy>=1.11 in /usr/local/lib/python3.10/dist-packages (from plotly.express) (1.26.4)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.20.0->plotly.express) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.20.0->plotly.express) (2024.2)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.20.0->plotly.express) (2024.1)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from patsy>=0.5->plotly.express) (1.16.0)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from plotly>=4.1.0->plotly.express) (9.0.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from plotly>=4.1.0->plotly.express) (24.1)

```
1 # Define the paths for train and test datasets in your Google Drive
2 train_path = '/content/drive/MyDrive/Colab Notebooks/train.csv'
3 test_path = '/content/drive/MyDrive/Colab Notebooks/test.csv'
4
5 # Load the datasets
6 import pandas as pd
7 train_df = pd.read_csv(train_path)
8 test_df = pd.read_csv(test_path)
```

```
1 # Check for missing values in train and test data
2 print(train_df.isnull().sum())
3
```

```
↗ ID          0
  Gender      0
  Ever_Married 140
  Age         0
  Graduated   78
  Profession  124
  Work_Experience 829
  Spending_Score 0
  Family_Size 335
  Var_1       76
  Segmentation 0
  dtype: int64
```



```
1 print(test_df.isnull().sum())
```

```
↗ ID          0
  Gender      0
  Ever_Married 50
  Age         0
  Graduated   24
  Profession   38
  Work_Experience 269
  Spending_Score 0
  Family_Size 113
  Var_1       32
  dtype: int64
```


```
1 !pip install missingno
2 !pip install plotly.express
3 import missingno as msno
4 import matplotlib.pyplot as plt
5 import plotly.express as px
6 # Visualize missing values
7 msno.matrix(train_df, figsize=(10, 6)) #removed extra indentation
8 plt.show() #removed extra indentation
9 msno.matrix(test_df, figsize=(10, 6)) #removed extra indentation
10 plt.show() #removed extra indentation
11
12 # Plot distributions of numerical columns
13 for column in train_df.select_dtypes(include=['float64', 'int64']).columns:
14     fig = px.histogram(train_df, x=column, title=f'Distribution of {column} in Train data')
15     fig.show()
16 for column in test_df.select_dtypes(include=['float64', 'int64']).columns:
17     fig = px.histogram(test_df, x=column, title=f'Distribution of {column} in Test data')
18     fig.show()
```



Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from missingno) (1.26.4)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from missingno) (3.7.1)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from missingno) (1.13.1)
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (from missingno) (0.13.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (1.3.0)
Requirement already satisfied:ycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (4.53.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (1.4.7)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (24.1)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (3.1.4)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->missingno) (2.8.2)
Requirement already satisfied: pandas>=1.2 in /usr/local/lib/python3.10/dist-packages (from seaborn->missingno) (2.1.4)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.2->seaborn->missingno) (2024.2)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.2->seaborn->missingno) (2024.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib->missingno) (1.16.0)
Requirement already satisfied: plotly.express in /usr/local/lib/python3.10/dist-packages (0.4.1)
Requirement already satisfied: pandas>=0.20.0 in /usr/local/lib/python3.10/dist-packages (from plotly.express) (2.1.4)
Requirement already satisfied: plotly>=4.1.0 in /usr/local/lib/python3.10/dist-packages (from plotly.express) (5.24.1)
Requirement already satisfied: statsmodels>=0.9.0 in /usr/local/lib/python3.10/dist-packages (from plotly.express) (0.14.3)
Requirement already satisfied: scipy>=0.18 in /usr/local/lib/python3.10/dist-packages (from plotly.express) (1.13.1)
Requirement already satisfied: patsy>=0.5 in /usr/local/lib/python3.10/dist-packages (from plotly.express) (0.5.6)
Requirement already satisfied: numpy>=1.11 in /usr/local/lib/python3.10/dist-packages (from plotly.express) (1.26.4)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.20.0->plotly.express) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.20.0->plotly.express) (2024.2)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.20.0->plotly.express) (2024.2)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from patsy>=0.5->plotly.express) (1.16.0)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from plotly>=4.1.0->plotly.express) (9.0.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from plotly>=4.1.0->plotly.express) (24.1)



2627




10



```
1 # Copy the original datasets to new variables
2 train_df_copy = train_df.copy()
3 test_df_copy = test_df.copy()
4
5 # Handle missing values for categorical variables (mode)
6 for df in [train_df_copy, test_df_copy]:
7     categorical_columns = df.select_dtypes(include=['object']).columns
8     for column in categorical_columns:
9         df[column].fillna(df[column].mode()[0], inplace=True)
10
11 # Handle missing values for numerical variables (median)
12 for df in [train_df_copy, test_df_copy]:
13     numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns
14     for column in numerical_columns:
15         df[column].fillna(df[column].median(), inplace=True)
16
17 # Visualize missing values for both train and test datasets
18 # Train dataset visualization
19 msno.matrix(train_df_copy, figsize=(10, 6))
20 plt.title('Missing Values Matrix After Filling in Train Data')
21 plt.show()
22
23 # Test dataset visualization
24 msno.matrix(test_df_copy, figsize=(10, 6))
25 plt.title('Missing Values Matrix After Filling in Test Data')
26 plt.show()
27
```




Missing Values Matrix After Filling in Train Data



Missing Values Matrix After Filling in Test Data



Perform encoding on categorical variables and scaling on numerical features and Prepare the Data

For train.csv:

- Split the data into X_train (features) and y_train (target labels).

For test.csv:

- Prepare the features (X_test), ensuring that all preprocessing steps (scaling, encoding) applied to train.csv are replicated on test.csv.

```
1 categorical_columns = X_train.select_dtypes(include=['object']).columns # Categorical columns
2 numerical_columns = X_train.select_dtypes(include=['float64', 'int64']).columns # Numerical columns
3
4 #Create a preprocessor that applies OneHotEncoder to categorical columns and StandardScaler to numerical columns
5 preprocessor = ColumnTransformer(
6     transformers=[
7         ('num', Pipeline(steps=[
8             ('imputer', SimpleImputer(strategy='median')), # Impute missing numerical data with median
9             ('scaler', StandardScaler()) # Scale numerical data
10        ]), numerical_columns),
11
12        ('cat', Pipeline(steps=[
13            ('imputer', SimpleImputer(strategy='most_frequent')), # Impute missing categorical data with most frequent
14            ('encoder', OneHotEncoder(drop='first')) # OneHotEncode categorical data
15        ]), categorical_columns)
16    ])
17
18 #Apply the preprocessor to X_train and fit the transformer
19 X_train_prepared = preprocessor.fit_transform(X_train)
20
21 #Apply the same preprocessor to the test data (X_test)
22 X_test_prepared = preprocessor.transform(test_data) # Only transform, do not fit on test data
23
24 #Check the shapes of the prepared data
25 print(f'Shape of X_train_prepared: {X_train_prepared.shape}')
26 print(f'Shape of X_test_prepared: {X_test_prepared.shape}')
27
```

```
➦ Shape of X_train_prepared: (8068, 23)
  Shape of X_test_prepared: (2627, 23)
```

```
1 # One-hot encode the target labels
2 y_train_transformed = pd.get_dummies(y_train).values
3
4 # Display the shape of the encoded labels
5 print("Shape of y_train_transformed:", y_train_transformed.shape)
```

```
➦ Shape of y_train_transformed: (8068, 4)
```

Model 1 - Best Practices Model

Build a neural network model with at least two hidden layers using the ReLU activation function.

Use softmax as the activation function for the output layer to handle multi-class classification.

Use categorical cross-entropy as the loss function and Adam optimizer.

```
1 model = Sequential([
2     Dense(64, activation='relu', input_shape=(X_train_prepared.shape[1],)), # First hidden layer
3     Dense(32, activation='relu'), # Second hidden layer
4     Dense(y_train_transformed.shape[1], activation='softmax') # Output layer
5 ])
6
7 # Step 4: Compile the model
8 model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
➦ /usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning:
```

Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the

Model Training:

Train the model using the train.csv data (X_train, y_train) for 50 epochs with a batch size of 32.

Use early stopping to prevent overfitting (monitor the validation loss and set a patience of 5 epochs).

Use 20% of the training data as a validation set during training

```

1 # Import necessary libraries
2 from tensorflow.keras.callbacks import EarlyStopping
3
4 # Step 1: Define early stopping callback
5 early_stopping = EarlyStopping(
6     monitor='val_loss',
7     patience=5,
8     restore_best_weights=True
9 )
10
11 # Step 2: Train the model with early stopping
12 history = model.fit(
13     X_train_prepared,
14     y_train_transformed,
15     epochs=50,
16     batch_size=32,
17     validation_split=0.2,
18     callbacks=[early_stopping],
19     verbose=1
20 )
21
22

```

```

Epoch 1/50
202/202 ————— 6s 5ms/step - accuracy: 0.3914 - loss: 1.2837 - val_accuracy: 0.5031 - val_loss: 1.1336
Epoch 2/50
202/202 ————— 1s 4ms/step - accuracy: 0.5061 - loss: 1.1118 - val_accuracy: 0.4957 - val_loss: 1.0961
Epoch 3/50
202/202 ————— 1s 5ms/step - accuracy: 0.5248 - loss: 1.0896 - val_accuracy: 0.5161 - val_loss: 1.0869
Epoch 4/50
202/202 ————— 1s 4ms/step - accuracy: 0.5352 - loss: 1.0566 - val_accuracy: 0.5130 - val_loss: 1.0804
Epoch 5/50
202/202 ————— 1s 4ms/step - accuracy: 0.5251 - loss: 1.0611 - val_accuracy: 0.5130 - val_loss: 1.0786
Epoch 6/50
202/202 ————— 1s 4ms/step - accuracy: 0.5432 - loss: 1.0429 - val_accuracy: 0.5124 - val_loss: 1.0747
Epoch 7/50
202/202 ————— 2s 5ms/step - accuracy: 0.5439 - loss: 1.0387 - val_accuracy: 0.5124 - val_loss: 1.0698
Epoch 8/50
202/202 ————— 1s 5ms/step - accuracy: 0.5515 - loss: 1.0198 - val_accuracy: 0.5217 - val_loss: 1.0721
Epoch 9/50
202/202 ————— 2s 8ms/step - accuracy: 0.5580 - loss: 1.0187 - val_accuracy: 0.5198 - val_loss: 1.0715
Epoch 10/50
202/202 ————— 3s 9ms/step - accuracy: 0.5566 - loss: 1.0214 - val_accuracy: 0.5186 - val_loss: 1.0662
Epoch 11/50
202/202 ————— 1s 3ms/step - accuracy: 0.5519 - loss: 1.0119 - val_accuracy: 0.5118 - val_loss: 1.0661
Epoch 12/50
202/202 ————— 1s 5ms/step - accuracy: 0.5685 - loss: 1.0028 - val_accuracy: 0.5211 - val_loss: 1.0674
Epoch 13/50
202/202 ————— 1s 4ms/step - accuracy: 0.5617 - loss: 1.0173 - val_accuracy: 0.5242 - val_loss: 1.0690
Epoch 14/50
202/202 ————— 1s 4ms/step - accuracy: 0.5587 - loss: 1.0040 - val_accuracy: 0.5198 - val_loss: 1.0640
Epoch 15/50
202/202 ————— 1s 4ms/step - accuracy: 0.5667 - loss: 0.9941 - val_accuracy: 0.5266 - val_loss: 1.0709
Epoch 16/50
202/202 ————— 1s 3ms/step - accuracy: 0.5774 - loss: 0.9750 - val_accuracy: 0.5310 - val_loss: 1.0711
Epoch 17/50
202/202 ————— 2s 6ms/step - accuracy: 0.5706 - loss: 0.9950 - val_accuracy: 0.5242 - val_loss: 1.0798
Epoch 18/50
202/202 ————— 1s 5ms/step - accuracy: 0.5674 - loss: 0.9838 - val_accuracy: 0.5149 - val_loss: 1.0694
Epoch 19/50
202/202 ————— 2s 7ms/step - accuracy: 0.5797 - loss: 0.9561 - val_accuracy: 0.5149 - val_loss: 1.0688

```

Model Summary:

Print the model architecture summary to review the structure.

```

1 # Print the model architecture summary
2 model.summary()
3 # Show the model structure
4 plot_model(model, to_file='model_1_structure.png', show_shapes=True, show_layer_names=True)
5 from IPython.display import Image
6 Image(filename='model_1_structure.png')

```