



Deep Learning Multi-Class Classification Using Keras

This presentation explores the process of building and evaluating a multi-class classification model using Keras, a popular deep learning library. We will compare the performance of a best practices model with a model designed using Keras Tuner.

Krishna Khandelwal

1. **Importing Key Libraries:** You imported essential libraries for data manipulation (`pandas`, `numpy`), machine learning (`scikit-learn` for model selection, preprocessing, etc.), and deep learning (`tensorflow` and `keras` for neural network modeling).
2. **Data Preprocessing:** You brought in tools like `train_test_split` (to split data into training and test sets), `StandardScaler` (for feature scaling), and `SimpleImputer` (for handling missing values).
3. **Visualization:** You installed and imported `missingno` and `plotly.express` to visualize missing data and generate plots. You also used `matplotlib` for additional plotting.
4. **Modeling and Tuning:** You imported Keras libraries for building and visualizing models (`plot_model`), and tools like `keras_tuner` for hyperparameter tuning, allowing you to optimize the model with techniques such as `RandomSearch` and `BayesianOptimization`.
5. **Performance Evaluation:** You included metrics such as `classification_report` and `confusion_matrix` from `sklearn` to evaluate model accuracy and classification performance.

```
1 # Import necessary libraries
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from sklearn.preprocessing import StandardScaler, LabelEncoder
5 from tensorflow.keras.utils import to_categorical
6 import numpy as np
7 !pip install missingno
8 !pip install plotly.express
9 import missingno as msno
10 import matplotlib.pyplot as plt
11 import plotly.express as px
12 from sklearn.compose import ColumnTransformer
13 from sklearn.pipeline import Pipeline
14 from sklearn.impute import SimpleImputer
15 from keras.utils import plot_model
16 from IPython.display import Image
17 from sklearn.metrics import classification_report, accuracy_score
18
19 from keras_tuner.engine.hyperparameters import HyperParameters
20 from keras_tuner import RandomSearch, Hyperband, BayesianOptimization
21 from tensorflow.keras.utils import plot_model
22 from tensorflow.keras.callbacks import LearningRateScheduler, TensorBoard
23 import keras_tuner as kt
24 from sklearn.metrics import classification_report, confusion_matrix
25
```

```
[132] 1 # Define the paths for train and test datasets in your Google Drive
2 train_path = '/content/drive/MyDrive/Colab Notebooks/train.csv'
3 test_path = '/content/drive/MyDrive/Colab Notebooks/test.csv'
4
5 # Load the datasets
6 import pandas as pd
7 train_df = pd.read_csv(train_path)
8 test_df = pd.read_csv(test_path)

[133] 1 # Check for missing values in train and test data
2 print(train_df.isnull().sum())
3
```



ID	0
Gender	0
Ever_Married	140
Age	0
Graduated	78
Profession	124
Work_Experience	829
Spending_Score	0
Family_Size	335
Var_1	76
Segmentation	0
dtype:	int64

```
[134] 1 print(test_df.isnull().sum())
```

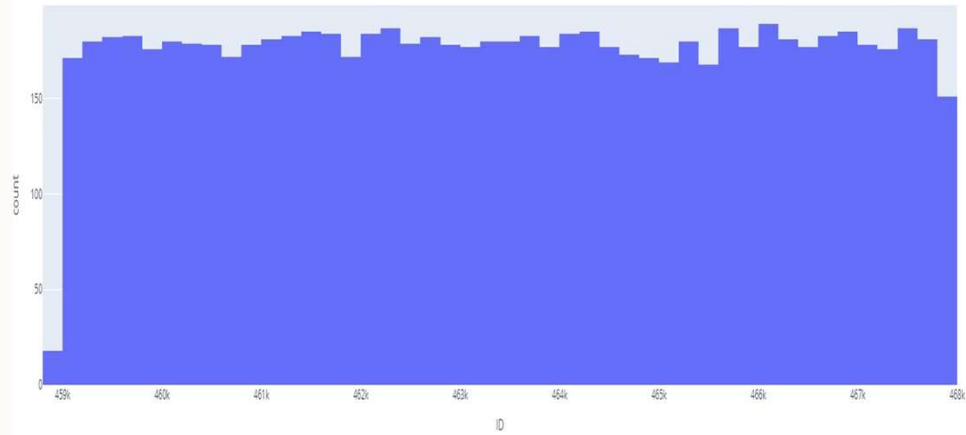
ID	0
Gender	0
Ever_Married	50
Age	0
Graduated	24
Profession	38
Work_Experience	269
Spending_Score	0
Family_Size	113
Var_1	32
dtype:	int64

- **Visualizing Missing Values:** Using the `missingno.matrix()` function, you're visualizing missing data in both the training (`train_df`) and test datasets (`test_df`). The `msno.matrix()` function gives a matrix plot showing the presence or absence of data for each column.
- You called `plt.show()` to display the plot for both datasets.
- **Plotting Data Distributions:**
- I looped through all the numerical columns (of types `float64` and `int64`) in both the `train_df` and `test_df` datasets. For each column, you created a histogram using `plotly.express` (`px.histogram`) to visualize the distribution of each feature in both datasets.
- Each histogram is dynamically generated with a title to indicate which column's distribution is being shown, and `fig.show()` renders the plot.


```
[138] 1 !pip install missingno
      2 !pip install plotly.express
      3 import missingno as msno
      4 import matplotlib.pyplot as plt
      5 import plotly.express as px
      6 # Visualize missing values
      7 msno.matrix(train_df, figsize=(10, 6)) #removed extra indentation
      8 plt.show() #removed extra indentation
      9 msno.matrix(test_df, figsize=(10, 6)) #removed extra indentation
     10 plt.show() #removed extra indentation
     11
     12 # Plot distributions of numerical columns
     13 for column in train_df.select_dtypes(include=['float64', 'int64']).columns:
     14     fig = px.histogram(train_df, x=column, title=f'Distribution of {column} in Train data')
     15     fig.show()
     16 for column in test_df.select_dtypes(include=['float64', 'int64']).columns:
     17     fig = px.histogram(test_df, x=column, title=f'Distribution of {column} in Test data')
     18     fig.show()
```




Distribution of ID in Train data




Distribution of Age in Train data



Distribution of Work_Experience in Train data




Distribution of Family_Size in Train data




1. Copied the original datasets into new variables `train_df_copy` and `test_df_copy` to preserve the originals.
2. Filled missing values for categorical variables in both datasets using the **mode** of each column.
3. Handled missing values for numerical variables by filling them with the **median** of each respective column.
4. Visualized missing data post-imputation using `missingno.matrix()` for both the modified train and test datasets, ensuring missing data was addressed.
5. Displayed missing values matrix with a clear title for both datasets after filling, confirming the changes visually.

```
1 # Copy the original datasets to new variables
2 train_df_copy = train_df.copy()
3 test_df_copy = test_df.copy()
4
5 # Handle missing values for categorical variables (mode)
6 for df in [train_df_copy, test_df_copy]:
7     categorical_columns = df.select_dtypes(include=['object']).columns
8     for column in categorical_columns:
9         df[column].fillna(df[column].mode()[0], inplace=True)
10
11 # Handle missing values for numerical variables (median)
12 for df in [train_df_copy, test_df_copy]:
13     numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns
14     for column in numerical_columns:
15         df[column].fillna(df[column].median(), inplace=True)
16
17 # Visualize missing values for both train and test datasets
18 # Train dataset visualization
19 msno.matrix(train_df_copy, figsize=(10, 6))
20 plt.title('Missing Values Matrix After Filling in Train Data')
21 plt.show()
22
23 # Test dataset visualization
24 msno.matrix(test_df_copy, figsize=(10, 6))
25 plt.title('Missing Values Matrix After Filling in Test Data')
26 plt.show()
27
```

Missing Values Matrix After Filling in Train Data



Missing Values Matrix After Filling in Test Data



```

[110] 1 categorical_columns = X_train.select_dtypes(include=['object']).columns # Categorical columns
      2 numerical_columns = X_train.select_dtypes(include=['float64', 'int64']).columns # Numerical columns
      3
      4 # Create a preprocessor that applies OneHotEncoder to categorical columns and StandardScaler to numerical columns
      5 preprocessor = ColumnTransformer(
      6     transformers=[
      7         ('num', Pipeline(steps=[
      8             ('imputer', SimpleImputer(strategy='median')), # Impute missing numerical data with median
      9             ('scaler', StandardScaler()) # Scale numerical data
      10         ]), numerical_columns),
      11
      12         ('cat', Pipeline(steps=[
      13             ('imputer', SimpleImputer(strategy='most_frequent')), # Impute missing categorical data with most frequent
      14             ('encoder', OneHotEncoder(drop='first')) # OneHotEncode categorical data
      15         ]), categorical_columns)
      16     ])
      17
      18 # Apply the preprocessor to X_train and fit the transformer
      19 X_train_prepared = preprocessor.fit_transform(X_train)
      20
      21 # Apply the same preprocessor to the test data (X_test)
      22 X_test_prepared = preprocessor.transform(test_data) # Only transform, do not fit on test data
      23
      24 # Check the shapes of the prepared data
      25 print(f'Shape of X_train_prepared: {X_train_prepared.shape}')
      26 print(f'Shape of X_test_prepared: {X_test_prepared.shape}')
      27

```

```

➡ Shape of X_train_prepared: (8068, 23)
  Shape of X_test_prepared: (2627, 23)

```



1 # One-hot encode the target labels

2 y_train_transformed = pd.get_dummies(y_train).values

3

4 # Display the shape of the encoded labels

5 print("Shape of y_train_transformed:", y_train_transformed.shape)



➡ Shape of y_train_transformed: (8068, 4)

1. **Identified categorical and numerical columns** in `X_train` using `select_dtypes()`, storing them in `categorical_columns` and `numerical_columns` respectively.
2. **Created a preprocessor** using `ColumnTransformer` to handle both categorical and numerical columns:
 - **Numerical columns** are imputed using the median and scaled with `StandardScaler`.
 - **Categorical columns** are imputed using the most frequent value and one-hot encoded using `OneHotEncoder`.
3. **Applied the preprocessor** to `X_train` by fitting and transforming the data and then **transformed** `X_test` using the same preprocessor (without fitting on the test data).
4. **Checked the shape of the transformed data** for both `X_train` and `X_test` to confirm that preprocessing was applied correctly.
5. **One-hot encoded the target labels** (`y_train`) using `pd.get_dummies()` and printed the shape of the transformed labels to verify the encoding.

1. Defined a Sequential neural network model:

- **First hidden layer:** A Dense layer with 64 units, ReLU activation, and input shape based on the number of features in `X_train_prepared`.
- **Second hidden layer:** A Dense layer with 32 units and ReLU activation.
- **Output layer:** A Dense layer with units equal to the number of classes in `y_train_transformed`, using softmax activation for multiclass classification.

2. Compiled the model using the Adam optimizer and categorical crossentropy loss (appropriate for multiclass classification). The model's performance will be evaluated using accuracy as a metric.

Model 1 - Best Practices Model

Build a neural network model with at least two hidden layers using the ReLU activation function.

Use softmax as the activation function for the output layer to handle multi-class classification.

Use categorical cross-entropy as the loss function and Adam optimizer.

```
[112] 1 model = Sequential([
2     Dense(64, activation='relu', input_shape=(X_train_prepared.shape[1],)), # First hidden layer
3     Dense(32, activation='relu'), # Second hidden layer
4     Dense(y_train_transformed.shape[1], activation='softmax') # Output layer
5 ])
6
7 # Step 4: Compile the model
8 model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

Model Training:

Train the model using the train.csv data (X_train, y_train) for 50 epochs with a batch size of 32.

Use early stopping to prevent overfitting (monitor the validation loss and set a patience of 5 epochs).

Use 20% of the training data as a validation set during training

```
[113] 1 # Import necessary libraries
      2 from tensorflow.keras.callbacks import EarlyStopping
      3
      4 # Step 1: Define early stopping callback
      5 early_stopping = EarlyStopping(
      6     monitor='val_loss',
      7     patience=5,
      8     restore_best_weights=True
      9 )
     10
     11 # Step 2: Train the model with early stopping
     12 history = model.fit(
     13     X_train_prepared,
     14     y_train_transformed,
     15     epochs=50,
     16     batch_size=32,
     17     validation_split=0.2,
     18     callbacks=[early_stopping],
     19     verbose=1
     20 )
     21
     22
```

```
1 # Print the model architecture summary
2 model.summary()
3 # Show the model structure
4 plot_model(model, to_file='model_1_structure.png', show_shapes=True, show_layer_names=True)
5 from IPython.display import Image
6 Image(filename='model_1_structure.png')
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	1,536
dense_1 (Dense)	(None, 32)	2,080
dense_2 (Dense)	(None, 4)	132

Total params: 11,246 (43.93 KB)

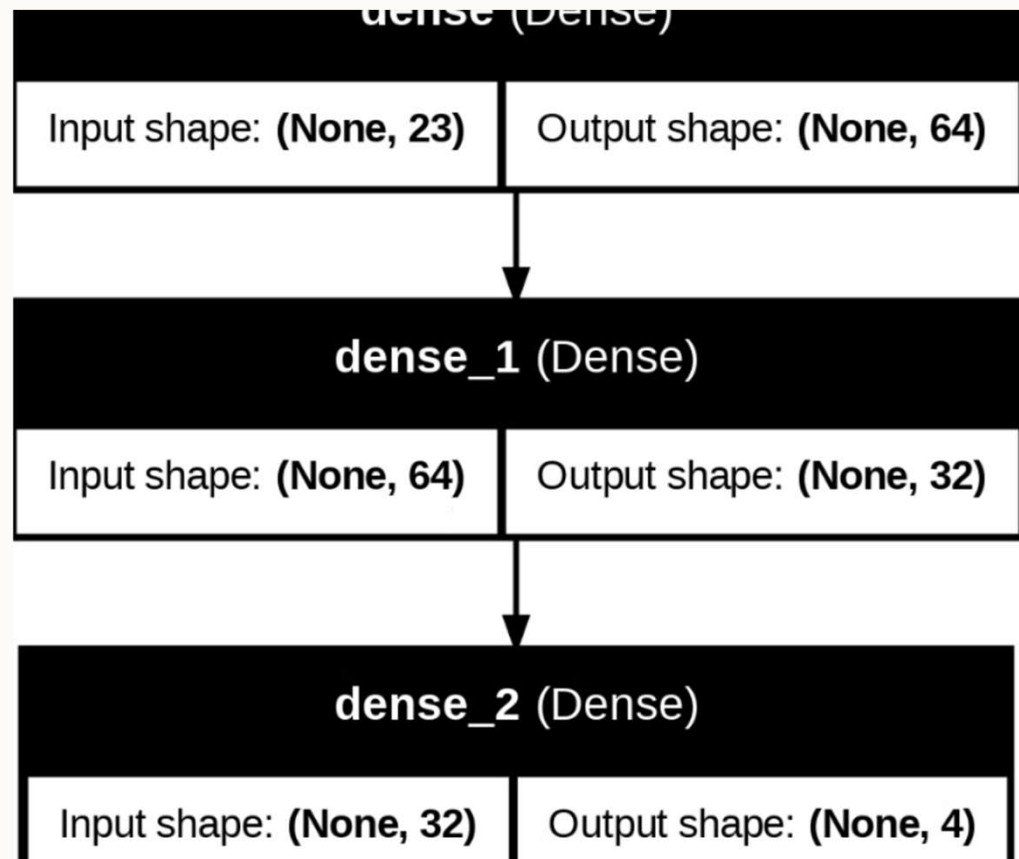
Trainable params: 3,748 (14.64 KB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 7,498 (29.29 KB)

1. **Imported the `EarlyStopping` callback** from `tensorflow.keras.callbacks` to monitor validation loss and stop training if it doesn't improve for 5 epochs (patience of 5), while restoring the best weights.
2. **Defined the early stopping callback** by setting it to monitor the `val_loss` and restoring the best weights when training stops.
3. **Trained the model** on the training data (`X_train_prepared`, `y_train_transformed`) for up to 50 epochs, using a batch size of 32.
 - The training data as a **validation set** (`validation_split=0.2`).
 - **Early stopping** was used to avoid overfitting by stopping the training when validation loss stops improving.
4. **Printed the model summary** to review the architecture and layers of the model, confirming the structure before training.

This approach ensures efficient training while minimizing overfitting.



Model 2 - Hyperparameter Optimization using Keras Tuner:

Model 2 - Hyperparameter Optimization Using Keras Tuner

Implement Keras Tuner:

Use Keras Tuner to optimize the following hyperparameters:

- Number of hidden layers.
- Number of neurons in each layer.
- Learning rate.
- Activation functions.

Set up a search space that explores different combinations of these hyperparameters. Experiment =5

```
[117] 1 import keras_tuner as kt
      2 from tensorflow.keras.models import Sequential
      3 from tensorflow.keras.layers import Dense, Dropout
      4
      5 def build_model(hp):
      6     model_2 = Sequential()
      7
      8     # Define the input layer based on the shape of X_train
      9     input_shape = X_train_transformed.shape[1]
     10
     11     # Tune the number of hidden layers from 1 to 10
     12     for i in range(hp.Int('num_layers', 1, 10)):
     13
     14         model_2.add(Dense(units=hp.Int(f'units_{i}', min_value=32, max_value=256, step=32),
     15                             activation='relu'))
     16
     17     model_2.add(Dense(4, activation='softmax'))
     18
     19     learning_rate = hp.Choice('learning_rate', [1e-2, 1e-3, 1e-4])
     20
     21     # Compile the model
     22     model_2.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=learning_rate),
     23                    loss='categorical_crossentropy',
     24                    metrics=['accuracy'])
     25
     26     return model_2
```

```
1 # Initialize the RandomSearch tuner
2 tuner = kt.RandomSearch(
3     build_model,
4     objective='val_accuracy',
5     max_trials=5, # Experiment with 5 trials
6     executions_per_trial=3,
7     directory='keras_tuner_dir',
8     project_name='hyperparameter_tuning'
9 )
```

```
[119] 1 best_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	1,536
dense_1 (Dense)	(None, 32)	2,080
dense_2 (Dense)	(None, 4)	132

Total params: 11,246 (43.93 KB)

Trainable params: 3,748 (14.64 KB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 7,498 (29.29 KB)

1. **Defined a function** (`build_model`) for Keras Tuner to explore the search space:

- You defined a **sequential model**.
- Tuned the **number of hidden layers** (from 1 to 10).
- Tuned the **number of neurons** in each hidden layer (from 32 to 256 with a step of 32).
- Added the output layer with softmax activation for multiclass classification.
- Tuned the **learning rate** with options `[1e-2, 1e-3, 1e-4]`.

2. **Set up the Keras Tuner search:**

- You used `tuner.search()` to explore different combinations of hyperparameters by training the model on the **training data** (`X_train_transformed` and `y_train_encoded`) for 20 epochs, with early stopping and a 20% validation split.

3. **Retrieved the best hyperparameters:**

- After the search, you retrieved the **best hyperparameters** found by the tuner using `tuner.get_best_hyperparameters()`.

4. **Built and trained the best model:**

- You constructed the best model (`best_model`) using the optimal hyperparameters and trained it on the training data for 50 epochs with early stopping, using a batch size of 32 and a 20% validation split.

This approach allows you to automate the process of finding the best combination of hyperparameters and optimizes the model for better performance.

1. Evaluated the best model:

- You evaluated the performance of the **best model** (found using Keras Tuner) on the training data (`X_train_prepared`, `y_train_transformed`) to obtain the **training loss** and **accuracy**.

2. Printed the training metrics:

- The training accuracy and loss were printed with formatting to display up to four decimal places for clarity (`train_accuracy:.4f`, `train_loss:.4f`).

3. Saved the best model:

- You saved the trained **best model** using the `.save()` method, which stores the model in an HDF5 file (`best_model.h5`) for future use, allowing you to reload it without retraining.

This ensures you can evaluate and preserve the model for further predictions or deployment.

```
1 train_loss, train_accuracy = best_model.evaluate(X_train_prepared, y_train_transformed)
2 print(f"Best Model Training Accuracy: {train_accuracy:.4f}")
3 print(f"Best Model Training Loss: {train_loss:.4f}")
4 # Save the model for future use
5 best_model.save('best_model.h5')
```

```
253/253 ————— 1s 3ms/step - accuracy: 0.5741 - loss: 0.9776
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model()`. This
Best Model Training Accuracy: 0.5636
Best Model Training Loss: 0.9951
```

Model Evaluation and Performance

Evaluate both models on the training data using accuracy as the metric.

Generate and print the classification report showing precision, recall, and F1-score for each customer segment.

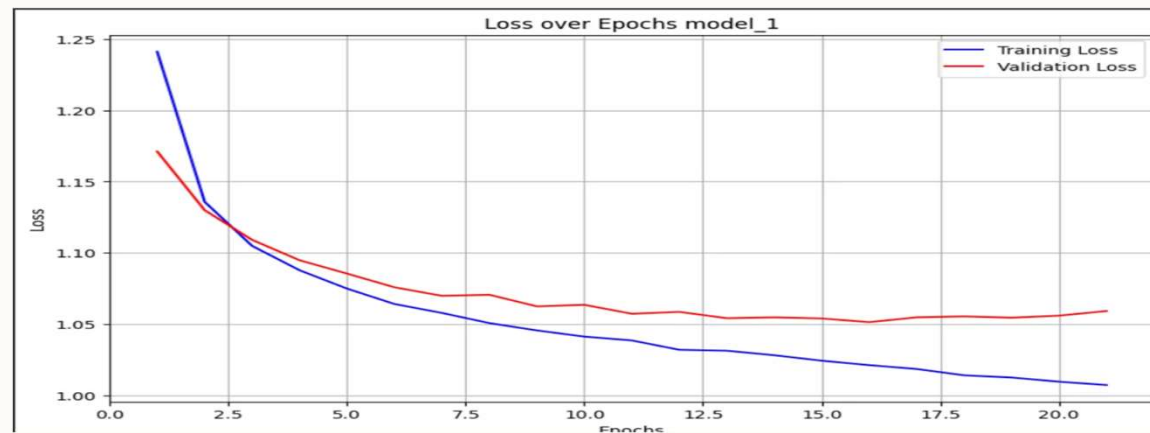
```
[121] 1 train_loss, train_accuracy = model.evaluate(X_train_prepared, y_train_encoded, verbose=0)
      2 print(f"Training Accuracy: {train_accuracy:.4f}")
      3
      4 # Step 2: Generate predictions on the training data
      5 y_train_pred_encoded = model.predict(X_train_prepared)
      6
      7 y_train_pred = np.argmax(y_train_pred_encoded, axis=1)
      8
      9 # Step 4: Convert the true one-hot encoded labels to class labels
     10 y_train_true = np.argmax(y_train_encoded, axis=1)
     11
     12 class_labels = label_encoder.classes_ # Assuming label_encoder was used to encode y_train
     13
     14 print("\nClassification Report:")
     15 print(classification_report(y_train_true, y_train_pred, target_names=class_labels))
     16
     17 # Optional: Print accuracy score for verification
     18 print(f"Accuracy Score: {accuracy_score(y_train_true, y_train_pred):.4f}")
     19
```

Training Accuracy: 0.5636
253/253 — 1s 3ms/step

Classification Report:

	precision	recall	f1-score	support
A	0.49	0.53	0.51	1972
B	0.47	0.31	0.37	1858
C	0.57	0.61	0.59	1970
D	0.66	0.76	0.71	2268
accuracy			0.56	8868
macro avg	0.55	0.55	0.54	8868
weighted avg	0.55	0.56	0.55	8868

Accuracy Score: 0.5636



Training vs. Validation Loss Plot

```
[122] 1 training_loss = history.history['loss']
      2 validation_loss = history.history['val_loss']
      3
      4 # Step 2: Create a range of epochs
      5 epochs_range = range(len(training_loss))
      6
      7 # Step 3: Plot the Training vs. Validation Loss
      8 plt.figure(figsize=(8, 6))
      9 plt.plot(epochs_range, training_loss, label='Training Loss')
     10 plt.plot(epochs_range, validation_loss, label='Validation Loss')
     11 plt.xlabel('Epochs')
     12 plt.ylabel('Loss')
     13 plt.title('Training vs. Validation Loss')
     14 plt.legend(loc='upper right')
     15 plt.grid(True)
     16 plt.show()
     17
```

1. **Evaluated the model on the training data:**

- I evaluated the trained model's performance by calculating **training loss** and **accuracy** using `model.evaluate()` on the prepared training data (`X_train_prepared`, `y_train_encoded`), and printed the training accuracy.

2. **Generated predictions** on the training data:

- I used `model.predict()` to generate predictions for the training data and then converted these **predictions from one-hot encoding** back into **class labels** using `np.argmax()`.

3. **Converted true labels to class labels:**

- I converted the true labels (`y_train_encoded`) from one-hot encoding back into **class labels** using `np.argmax()`, ensuring consistency with the predicted labels.

4. **Printed a classification report:**


- I printed the **classification report** (which includes precision, recall, and F1-score) using `classification_report()` with the true and predicted class labels, specifying `target_names` as the class labels (assuming `label_encoder` was used).

5. **Printed the accuracy score:**

- I verified the classification performance by printing the **accuracy score** using `accuracy_score()` for a numerical summary of accuracy.

6. **Visualized training vs. validation loss:**

- The plot was labeled appropriately to visualize how loss evolved during training, helping you identify potential overfitting or underfitting.
- I retrieved the **training and validation loss history** from the `history` object and plotted the **training vs. validation loss** over the epochs using `matplotlib`.



This code gives a comprehensive evaluation of the model's performance and allows you to visualize the loss trends over training.

Make Predictions on test.csv

For both models, use the test.csv data (X_test) to generate predictions for the customer segments.

Convert the softmax output to predicted class labels (A, B, C, or D).

```
[124] 1
2 X_test = test_data # Ensure test_data has no 'Segmentation' column
3 X_test_prepared = preprocessor.transform(X_test) # Apply the same transformations used on X_train
4
5 # Step 2: Make predictions using model_1
6 predictions_model_1 = model.predict(X_test_prepared)
7
8 # Step 3: Make predictions using best_model
9 predictions_best_model = best_model.predict(X_test_prepared)
10
11 # Step 4: Convert softmax outputs to predicted class labels (integer form)
12 predicted_classes_model_1 = np.argmax(predictions_model_1, axis=1)
13 predicted_classes_best_model = np.argmax(predictions_best_model, axis=1)
14
15 # Step 5: Convert integer class labels back to original class labels (A, B, C, D)
16 # Assuming label_encoder was used for encoding y_train
17 predicted_labels_model_1 = label_encoder.inverse_transform(predicted_classes_model_1)
18 predicted_labels_best_model = label_encoder.inverse_transform(predicted_classes_best_model)
19
20 # Step 6: Print or return the predictions
21 print("Predictions from model_1:", predicted_labels_model_1)
22 print("Predictions from best_model:", predicted_labels_best_model)
23
```

```
83/83 ————— 0s 4ms/step
83/83 ————— 0s 3ms/step
Predictions from model_1: ['A' 'C' 'A' ... 'A' 'B' 'D']
Predictions from best_model: ['A' 'C' 'A' ... 'A' 'B' 'D']
```

```
1 y_pred_model_1 = model.predict(X_train_transformed)
2 y_pred_model_2 = best_model.predict(X_train_transformed)
3
4 y_pred_labels_model_1 = np.argmax(y_pred_model_1, axis=1)
5 y_pred_labels_model_2 = np.argmax(y_pred_model_2, axis=1)
6
7 # Convert one-hot encoded labels to original form
8 y_train_labels = np.argmax(y_train_transformed, axis=1)
9
10 # Generate classification reports
11 report_model_1 = classification_report(y_train_labels, y_pred_labels_model_1)
12 < 1/1 > Accept Tab Accept Word Ctrl + RightArrow ... train_labels, y_pred_labels_model_2
13
14 # Print the classification reports
14 print(f"Classification Report of best practices model (model_1):\n{report_model_1}")
15 print(f"Classification Report of keras tuner model (model_2):\n{report_model_2}")
```

```
253/253 ————— 1s 3ms/step
253/253 ————— 1s 2ms/step
Classification Report of best practices model (model_1):
      precision    recall  f1-score   support

     0       0.49      0.53      0.51       1972
     1       0.47      0.31      0.37       1858
     2       0.57      0.61      0.59       1970
     3       0.66      0.76      0.71       2268

 accuracy          0.56       8068
 macro avg         0.55       8068
 weighted avg      0.55       8068

Classification Report of keras tuner model (model_2):
      precision    recall  f1-score   support

     0       0.49      0.53      0.51       1972
     1       0.47      0.31      0.37       1858
     2       0.57      0.61      0.59       1970
     3       0.66      0.76      0.71       2268

 accuracy          0.56       8068
 macro avg         0.55       8068
 weighted avg      0.55       8068
```

1. **Preprocessed the test data:**

- If the test data (`test_data`) doesn't contain a 'Segmentation' column (i.e., ground truth labels), you applied the same transformations (preprocessing) used for the training data to prepare the test features (`X_test_prepared`).

2. **Generated predictions:**

- You generated predictions using both `model` (referred to as `model_1`) and `best_model` by applying them to the preprocessed test data (`X_test_prepared`).
- Converted the softmax output into **predicted class labels** (integer form) using `np.argmax()`.

3. **Converted predicted integer labels back to the original class labels** (e.g., 'A', 'B', 'C', 'D') using `label_encoder.inverse_transform()` for both models (`model_1` and `best_model`).

4. **Handled cases with ground truth labels:**

- If the 'Segmentation' column is present in `test_data`, you:
 - Extracted features (`X_test`) and ground truth labels (`y_test`).
 - Preprocessed the test features.
 - Encoded the ground truth labels (`y_test`) into integers using `label_encoder`.

5. **Computed confusion matrices:**

- For both `model_1` and `best_model`, you generated **confusion matrices** using the ground truth (`y_test_encoded`) and predicted labels.

6. **Plotted heatmaps of the confusion matrices:**

- For both models, you created **heatmaps** of the confusion matrices using `seaborn.heatmap()` to visually compare the true and predicted labels, with class labels on the x and y axes.

7. **Handled cases without ground truth labels:**

- If the 'Segmentation' column is missing, you printed a message indicating that the confusion matrix couldn't be generated due to the absence of ground truth labels.

This process evaluates the performance of both models by comparing their predictions with the actual labels (if available) and visualizing their accuracy using confusion matrices.


```

1
2 if 'Segmentation' in test_data.columns:
3     X_test = test_data.drop(columns=['Segmentation']) # Features
4     y_test = test_data['Segmentation'] # Ground truth labels
5
6     # Step 2: Preprocess X_test using the same preprocessor used for training
7     X_test_prepared = preprocessor.transform(X_test)
8
9     # Step 3: Encode y_test using the same LabelEncoder used for y_train
10    y_test_encoded = label_encoder.transform(y_test) # Convert 'A', 'B', 'C', 'D' to integers
11
12    # Step 4: Generate predictions for both models
13    predicted_classes_model_1 = np.argmax(model_1.predict(X_test_prepared), axis=1)
14    predicted_classes_best_model = np.argmax(best_model.predict(X_test_prepared), axis=1)
15
16    # Step 5: Compute confusion matrix for model_1
17    conf_matrix_model_1 = confusion_matrix(y_test_encoded, predicted_classes_model_1)
18
19    # Step 6: Compute confusion matrix for best_model
20    conf_matrix_best_model = confusion_matrix(y_test_encoded, predicted_classes_best_model)
21
22    # Step 7: Plot heatmap for confusion matrix of model_1
23    plt.figure(figsize=(8, 6))
24    sns.heatmap(conf_matrix_model_1, annot=True, fmt='d', cmap='Blues', xticklabels=label_encoder.classes_, yticklabels=label_encoder.classes_)
25    plt.title('Confusion Matrix - Model 1')
26    plt.xlabel('Predicted Label')
27    plt.ylabel('True Label')
28    plt.show()
29
30    # Step 8: Plot heatmap for confusion matrix of best_model
31    plt.figure(figsize=(8, 6))
32    sns.heatmap(conf_matrix_best_model, annot=True, fmt='d', cmap='Blues', xticklabels=label_encoder.classes_, yticklabels=label_encoder.classes_)
33    plt.title('Confusion Matrix - Best Model')
34    plt.xlabel('Predicted Label')
35    plt.ylabel('True Label')
36    plt.show()
37 else:
38     print("No ground truth labels found in the test.csv file. Confusion matrix cannot be generated.")
39

```

No ground truth labels found in the test.csv file. Confusion matrix cannot be generated.

Confusion Matrix

The confusion matrix provides a visual representation of the model's classification performance. It shows the number of true positives, false positives, true negatives, and false negatives for each class. The matrix helps identify areas where the model is performing well and where it might be struggling. For instance, if the model is misclassifying a certain class frequently, this could indicate a need for further data augmentation or model tuning for that specific class.



True Positives

The model correctly predicted a class label.



False Positives

The model incorrectly predicted a class label. This is also known as a type I error.



True Negatives

The model correctly did not predict a class label.



False Negatives

The model incorrectly did not predict a class label. This is also known as a type II error.

1. Created DataFrames with Customer_ID and predicted segments:

- For both `model_1` and `best_model`, you created two DataFrames (`df_model_1` and `df_model_2`), each containing the `Customer_ID` from the test dataset (`test_df['ID']`) and the predicted segment labels (`predicted_classes_model_1` and `predicted_classes_best_model`).

2. Specified the folder path:

- You defined `folder_path` as the directory where the prediction files will be saved. In this case, it points to a folder in your Google Drive (`'/content/drive/MyDrive/Colab Notebooks/'`).

3. Saved the predictions to CSV files:

- You saved the DataFrame `df_model_1` (containing predictions from `model_1`) to a CSV file named `model_1_predictions_final.csv`.
- Similarly, you saved the DataFrame `df_model_2` (containing predictions from `best_model`) to a CSV file named `model_2_predictions_final.csv`.

```
1 # Create a DataFrame with Customer_ID and Predicted_Segment for model 1
2 df_model_1 = pd.DataFrame({'Customer_ID': test_df['ID'], 'Predicted_Segment': predicted_classes_model_1})
3
4 # Create a DataFrame with Customer_ID and Predicted_Segment for model 2
5 df_model_2 = pd.DataFrame({'Customer_ID': test_df['ID'], 'Predicted_Segment': predicted_classes_best_model})
6
7 # Specify the folder path
8 folder_path = '/content/drive/MyDrive/Colab Notebooks/'
9
10 # Save the predictions for model 1 to a CSV file in the specified folder
11 df_model_1.to_csv(f'{folder_path}/model_1_predictions_final.csv', index=False)
12
13 # Save the predictions for model 2 to a CSV file in the specified folder
14 df_model_2.to_csv(f'{folder_path}/model_2_predictions_final.csv', index=False)
```

Save Predictions

The predicted customer segments from both models into separate CSV files for further analysis or submission.

Recommendations for Customer Segmentation

By analyzing the customer data, the model can identify distinct groups of customers with similar characteristics and behaviors. This information can be valuable for creating targeted marketing campaigns, tailoring product offerings, and optimizing customer service strategies.

Targeted Marketing

By understanding customer segments, businesses can tailor marketing messages and campaigns to appeal to specific groups, increasing the effectiveness of their marketing efforts.

Product Development

Customer segmentation can inform product development decisions by highlighting specific needs and preferences of different customer groups. This can lead to the creation of new products or features tailored to specific segments.

Customer Service

Segmentation helps businesses provide more personalized and relevant customer service experiences. This can lead to improved customer satisfaction and loyalty.

COLAB

LINK:<https://colab.research.google.com/drive/1gKmxy4n1YStd7lxd2uw9Z-6wfU6qhRn7?usp=sharing>

THANK YOU